

HIVE CASE STUDY

SUBMITTED BY:

SUSHANT YADAV
NURSHAFIZAH MOHD KAMIL

PROBLEM STATEMENT:

On the clickstream cosmetic data, we are required to gain insights and analyse our customer browsing pattern which will help to increase profit of the company. This can be done through analysing customer browsing pattern on which products, purchases and views when they are browsing the tracking the clicks and pattern.

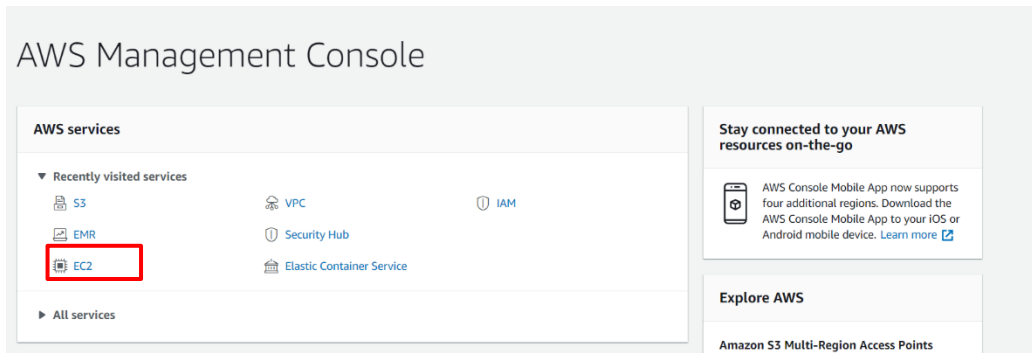
IMPLEMENTATION PHASE:

The implementation phase can be divided into the following parts:

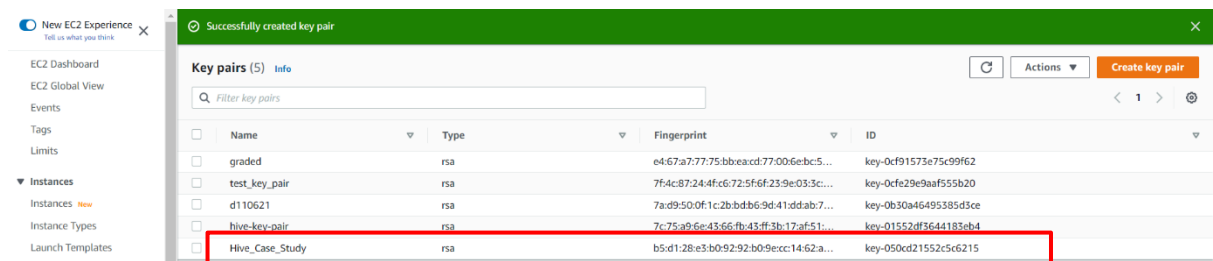
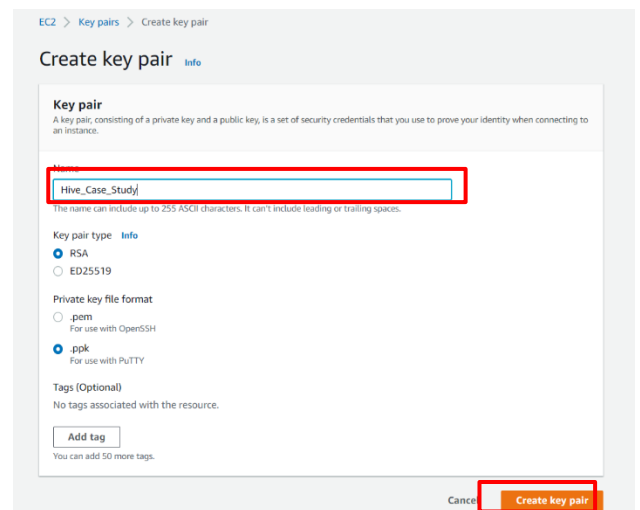
- Copying the data set into the HDFS:
- Launch an EMR cluster that utilizes the Hive services, and
- Move the data from the S3 bucket into the HDFS
- Creating the database and launching Hive queries on your EMR cluster:
- Create the structure of your database,
- Use optimized techniques to run your queries as efficiently as possible
- Show the improvement of the performance after using optimization on any single query.
- Run Hive queries to answer the questions given below.
- Cleaning up
- Drop your database, and
- Terminate your cluster

STEPS: CREATING KEY-PAIR

On the AWS Services, click on EC2 then 'Create key pair'.



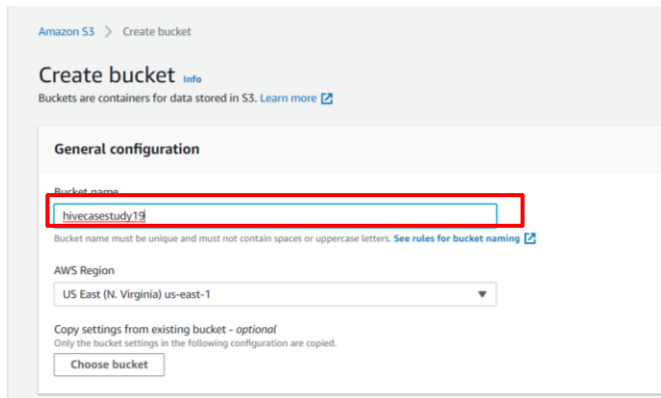
Fill the Name section, as in our case, our key pair name is 'Hive_Case_Study'. It is directly created as .ppk file. Next click on 'Create key pair'.



Our key pair 'Hive_Case_Study' has been created and downloaded.

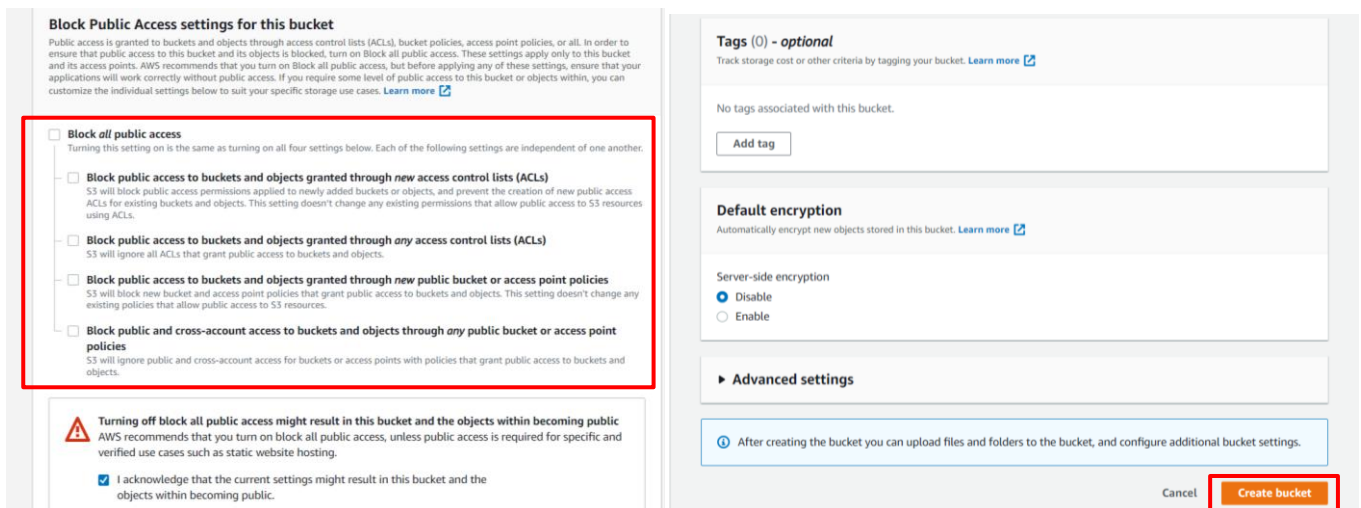
STEPS: CREATING BUCKET FOR THE CASE STUDY

We need to store our files in S3, thus the first step to serve the purpose is to 'Create Bucket'. On 'Amazon S3' go to 'Create bucket'.

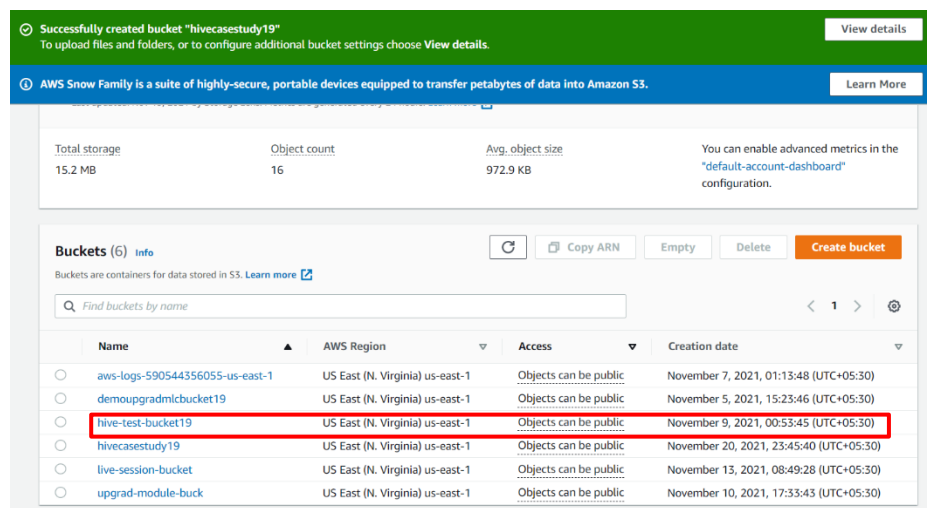


Field our Bucket name, do note that the bucket name should be unique.

When we are creating our Bucket, we need to enable our Public Access settings for the Bucket. Untick the following option > Create Bucket.

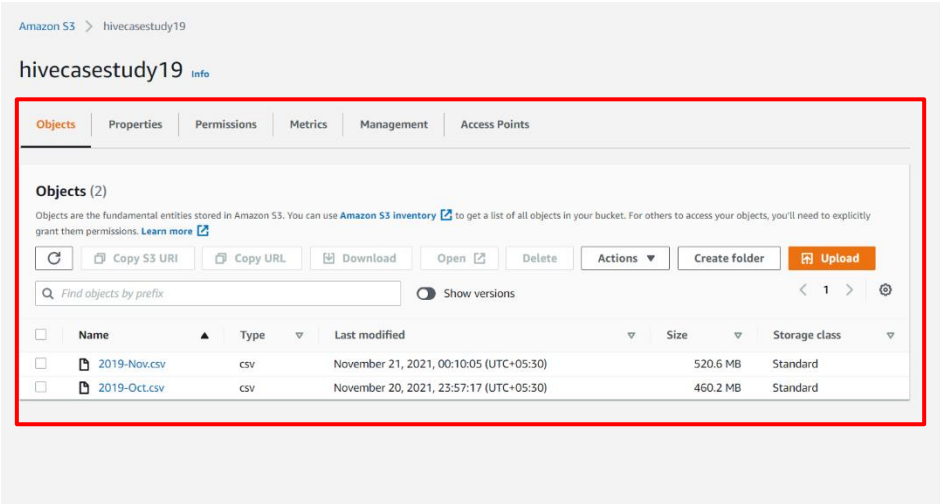


As we can see, our bucket 'hivecasestudy19' has successfully created. We should now be ready to upload our 2019-Oct.csv and 2019-Nov.csv files into the Bucket.



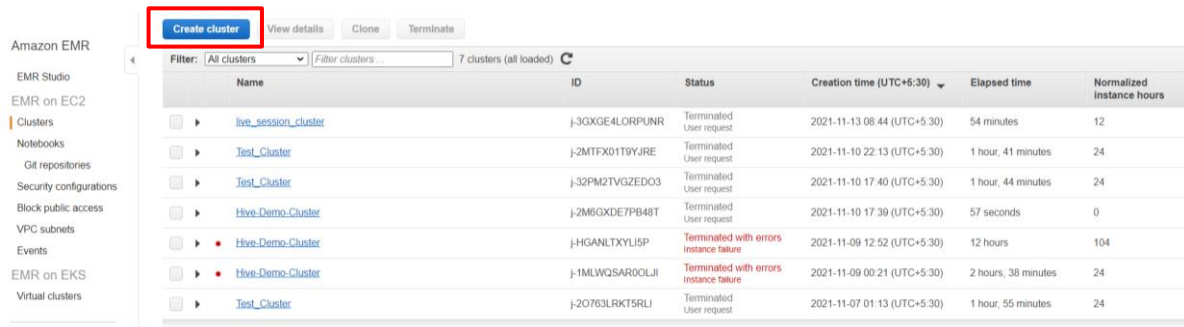
Name	AWS Region	Access	Creation date
aws-logs-590544356055-us-east-1	US East (N. Virginia) us-east-1	Objects can be public	November 7, 2021, 01:13:48 (UTC+05:30)
demougradmlcbucket19	US East (N. Virginia) us-east-1	Objects can be public	November 5, 2021, 15:23:46 (UTC+05:30)
hive-test-bucket19	US East (N. Virginia) us-east-1	Objects can be public	November 9, 2021, 00:53:45 (UTC+05:30)
hivecasestudy19	US East (N. Virginia) us-east-1	Objects can be public	November 20, 2021, 23:45:40 (UTC+05:30)
live-session-bucket	US East (N. Virginia) us-east-1	Objects can be public	November 13, 2021, 08:49:28 (UTC+05:30)
upgrad-module-buck	US East (N. Virginia) us-east-1	Objects can be public	November 10, 2021, 17:33:43 (UTC+05:30)

In our Bucket, we uploaded both of our files successfully. We will proceed now with our cluster creation.



STEPS: EMR (ELASTIC MAP REDUCE) – CREATING CLUSTER

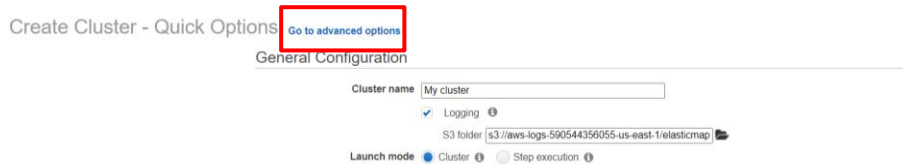
When we type EMR in the AWS Services, the below screen will pop up. As you can see, we are trying to create a cluster for this purpose of case study. Hence, below are the next steps.



The screenshot shows the Amazon EMR console. On the left sidebar, the 'Create cluster' button is highlighted with a red box. The main area displays a table of existing clusters. The table has columns for Name, ID, Status, Creation time (UTC+5:30), Elapsed time, and Normalized instance hours. The clusters listed include 'live_session_cluster', 'Test_Cluster', and 'Hive-Demo-Cluster'.

Name	ID	Status	Creation time (UTC+5:30)	Elapsed time	Normalized instance hours
live_session_cluster	j-3GXGE4LORPUNR	Terminated User request	2021-11-13 08:44 (UTC+5:30)	54 minutes	12
Test_Cluster	j-2MTFX01T9YJRE	Terminated User request	2021-11-10 22:13 (UTC+5:30)	1 hour, 41 minutes	24
Test_Cluster	j-32PM2TVGZED03	Terminated User request	2021-11-10 17:40 (UTC+5:30)	1 hour, 44 minutes	24
Hive-Demo-Cluster	j-2M8GXDE7PB48T	Terminated User request	2021-11-10 17:39 (UTC+5:30)	57 seconds	0
Hive-Demo-Cluster	j-HGANLTXYL5P	Terminated with errors Instance failure	2021-11-09 12:52 (UTC+5:30)	12 hours	104
Hive-Demo-Cluster	j-1MLWQSAR00LJI	Terminated with errors Instance failure	2021-11-09 00:21 (UTC+5:30)	2 hours, 38 minutes	24
Test_Cluster	j-20763LRKT5RLI	Terminated User request	2021-11-07 01:13 (UTC+5:30)	1 hour, 55 minutes	24

Click on the 'Create Cluster' then 'Go to advanced options'.



The screenshot shows the 'Create Cluster - Quick Options' screen. The 'Go to advanced options' button is highlighted with a red box. The screen displays the 'General Configuration' section with fields for Cluster name, Logging, S3 folder, and Launch mode.

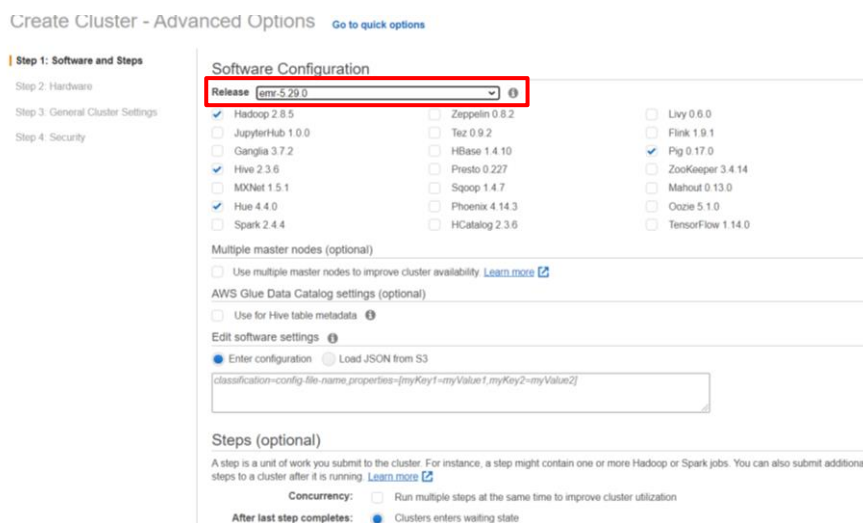
Cluster name: My cluster

Logging: ☒ Logging

S3 folder: s3://aws-logs-590544356055-us-east-1-elasticmap

Launch mode: ☒ Cluster ☐ Step execution

On the Advance Options, 'Step 1: Software and Steps' we will choose the **EMR RELEASED 5.29.0** for this case study. And we clicked 'Next'.



The screenshot shows the 'Create Cluster - Advanced Options' screen. The 'Release' dropdown menu is highlighted with a red box, showing 'emr-5.29.0'. The screen displays the 'Software Configuration' section with various software options and settings.

Release: **emr-5.29.0**

Software Configuration:

- ☒ Hadoop 2.8.5
- ☒ Hive 2.3.6
- ☒ Hue 4.4.0
- ☒ Spark 2.4.4
- ☐ Tez 0.9.2
- ☐ HBase 1.4.10
- ☐ Presto 0.227
- ☐ Sqoop 1.4.7
- ☐ Phoenix 4.14.3
- ☐ HCatalog 2.3.6
- ☐ Livy 0.6.0
- ☐ Flink 1.9.1
- ☒ Pig 0.17.0
- ☐ ZooKeeper 3.4.14
- ☐ Mahout 0.13.0
- ☐ Oozie 5.1.0
- ☐ TensorFlow 1.14.0

Multiple master nodes (optional): ☐ Use multiple master nodes to improve cluster availability. [Learn more](#)

AWS Glue Data Catalog settings (optional): ☐ Use for Hive table metadata

Edit software settings: ☒ Enter configuration ☐ Load JSON from S3

classification=config-file-name.properties-myKey1=myValue1-myKey2=myValue2

Steps (optional):

A step is a unit of work you submit to the cluster. For instance, a step might contain one or more Hadoop or Spark jobs. You can also submit additional steps to a cluster after it is running. [Learn more](#)

Concurrency: ☐ Run multiple steps at the same time to improve cluster utilization

After last step completes: ☒ Clusters enters waiting state

Note that for the purpose of case study, we have selected EMR Released 5.29.0 as some queries might run longer in the latest released.

Our cluster name in the case study is 'Hive Case Study'. And then, we clicked 'Next'

Create Cluster - Advanced Options [Go to quick options](#)

Step 1: Software and Steps
Step 2: Hardware
Step 3: General Cluster Settings
Step 4: Security

General Options

Cluster name: **Hive Case Study**

☒ Logging ⓘ
S3 folder: s3://aws-logs-590544356055-us-east-1/elasticmapreduce

☒ Debugging ⓘ
☒ Termination protection ⓘ

Tags ⓘ

Key	Value (optional)
Add a key to create a tag	

Additional Options

☐ EMRFS consistent view ⓘ
Custom AMI ID: **None** ⓘ

Bootstrap Actions

[Cancel](#) [Previous](#) [Next](#)

In the Security Options, we will be using our EC2 Key-Pair 'Hive_Case_Study' which had been downloaded earlier, and we clicked on 'Create Cluster'.

Create Cluster - Advanced Options [Go to quick options](#)

Step 1: Software and Steps
Step 2: Hardware
Step 3: General Cluster Settings
Step 4: Security

Security Options

EC2 key pair: **Hive_Case_Study** ⓘ

☒ Cluster visible to all IAM users in account ⓘ

Permissions ⓘ
☒ Default ☐ Custom
Use default IAM roles. If roles are not present, they will be automatically created for you with managed policies for automatic policy updates.

EMR role: [EMR_DefaultRole](#) ⓘ ☐ Use EMR_DefaultRole_V2 ⓘ

EC2 Instance profile: [EMR_EC2_DefaultRole](#) ⓘ

Auto Scaling role: [EMR_AutoScaling_DefaultRole](#) ⓘ

Security Configuration

EC2 security groups

[Cancel](#) [Previous](#) [Create cluster](#)

As you can see, our cluster is ready. For this case study, we will be using **2 node cluster** which consist of **1 Master Node (m4.large)** and **1 Core Node (m4.large)**. And both of these nodes are ready and running.

Amazon EMR

EMR Studio
EMR on EC2
Clusters
Notebooks
Git repositories
Security configurations
Block public access
VPC subnets
Events
EMR on EKS
Virtual clusters
Help
What's new

Cluster: Hive Case Study **Waiting** Cluster ready after last step completed.

Summary Application user interfaces Monitoring Hardware Configurations Events Steps Bootstrap actions

Summary

ID: j-2D00Q3PA8IFW5
Creation date: 2021-11-21 00:29 (UTC+5:30)
Elapsed time: 20 minutes
After last step completes: Cluster waits
Termination protection: On [Change](#)
Tags: [View All](#) / [Edit](#)
Master public DNS: ec2-54-80-51-78.compute-1.amazonaws.com [Connect to the Master Node Using SSH](#)

Configuration details

Release label: emr-5.29.0
Hadoop distribution: Amazon 2.8.5
Applications: Hive 2.3.6, Pig 0.17.0, Hue 4.4.0
Log URI: s3://aws-logs-590544356055-us-east-1/elasticmapreduce/ ⓘ
EMRFS consistent view: Disabled
Custom AMI ID: --

Application user interfaces

Persistent user interfaces ⓘ: --
On-cluster user Not Enabled [Enable an SSH Connection](#)
Interfaces ⓘ: --

Network and hardware

Availability zone: us-east-1a
Subnet ID: [subnet-2326527c](#) ⓘ
Master: Running 1 m4.large
Core: Running 1 m4.large
Task: --
Cluster scaling: Not enabled

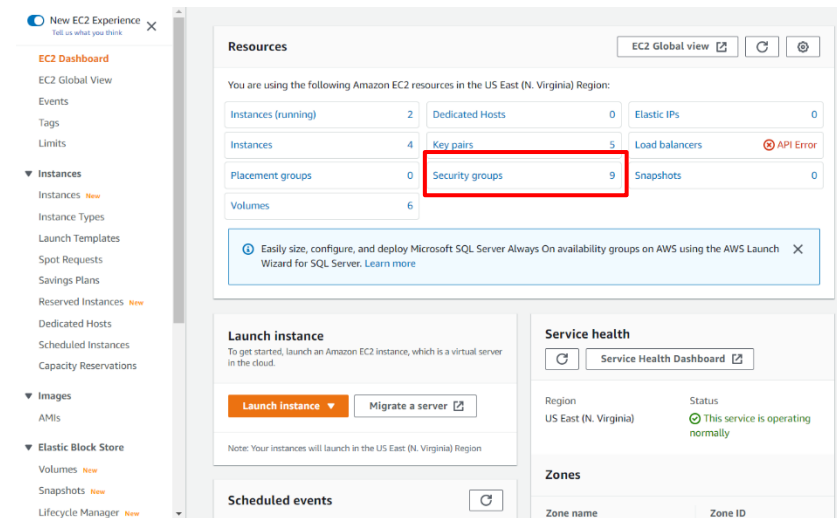
Security and access

Key name: Hive_Case_Study
EC2 Instance profile: EMR_EC2_DefaultRole
EMR role: EMR_DefaultRole
Auto Scaling role: EMR_AutoScaling_DefaultRole
Visible to all users: All [Change](#)
Security groups for Master: [sg-0e56b2be767c02e79](#) ⓘ (ElasticMapReduce-master)

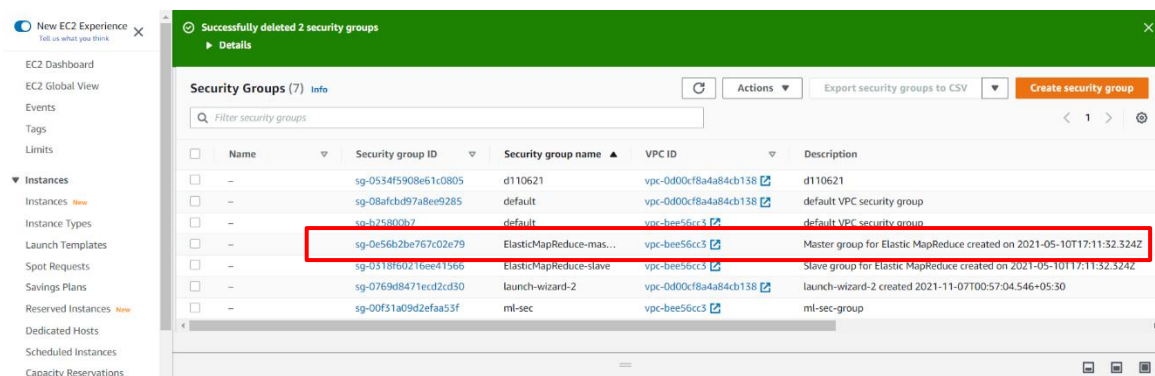
STEPS: SECURITY GROUPS

Once the cluster has successfully created, we have to edit the 'inbound rule'.


We will go to the EC2 and click on the 'Security Groups'.



Click on the Master Group.

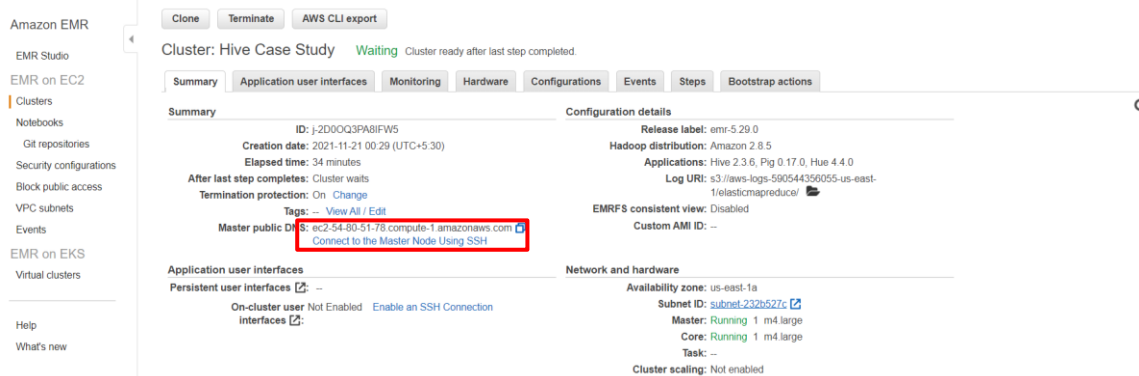


Save the 'SSH' inbound rules to 'Anywhere'

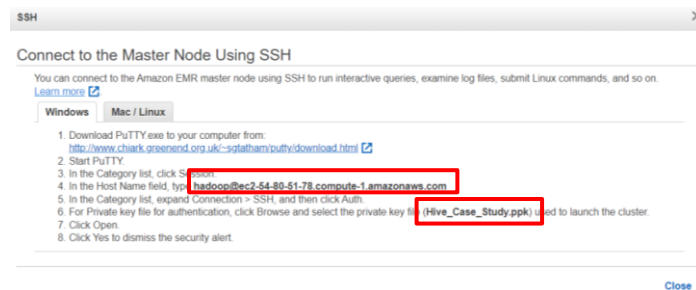


STEPS: CONNECTING THE MASTER NODE WITH PUTTY

Go to our ready cluster and click on the 'Connect to the Master Node Using SSH', another screen will pop up:



Important notes to be remembered here is our 'Host Name field' as well as our PPK file name saved earlier. Copied the 'Host Name field'.

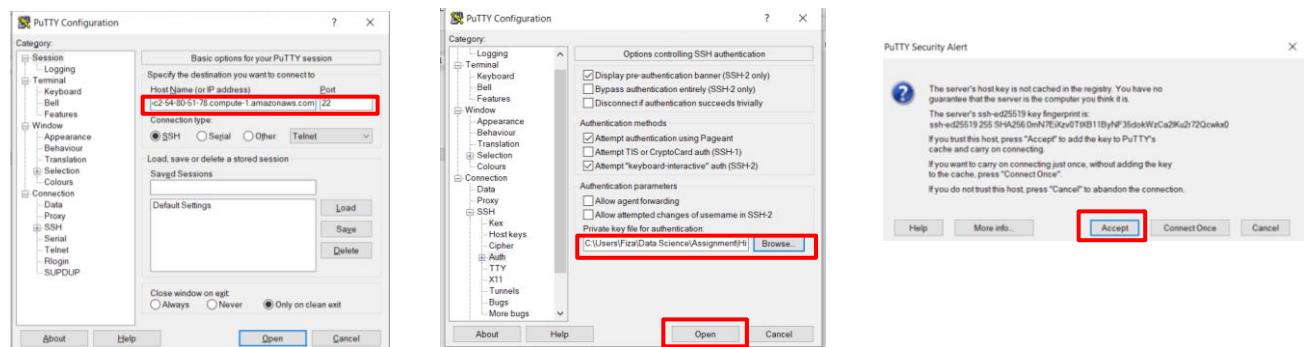


Host Name:


hadoop@ec2-54-80-51-78.compute-1.amazonaws.com

PPK file: Hive_Case_Study.ppk

Launch 'Putty', paste the 'Host Name' address which had been copied earlier in the 'Host Name' field. On the left side, click on 'SSH' > 'Auth' then browse the PPK file, in our case 'Hive_Case_Study.ppk', click 'Open' the 'Accept'.



EMR (Elastic Map Reduce) CLI has successfully launched.



STEPS: QUERYING IN HADOOP AND HIVE

Check the services running in Hadoop cluster

Command : ***sudo initctl list***

```
[hadoop@ip-172-31-44-156 ~]$ sudo initctl list
rc stop/waiting
tty (/dev/tty3) start/running, process 4946
tty (/dev/tty2) start/running, process 4944
tty (/dev/tty1) start/running, process 4941
tty (/dev/tty6) start/running, process 4957
tty (/dev/tty5) start/running, process 4952
tty (/dev/tty4) start/running, process 4948
update-motd stop/waiting
hive-server2 start/running, process 14480
hadoop-mapreduce-historyserver start/running, process 12867
hadoop-yarn-timelineserver start/running, process 12154
plymouth-shutdown stop/waiting
whisper-server stop/waiting
control-alt-delete stop/waiting
hive-hcatalog-server start/running, process 15245
```

Load the data sets into HDFS from S3

1. Verifying the inbuilt file directories in HDFS.

Command: ***hadoop fs -ls /***

```
[hadoop@ip-172-31-44-156 ~]$ hadoop fs -ls /
Found 4 items
drwxr-xr-x - hdfs hadoop 0 2021-11-20 19:05 /apps
drwxrwxrwt - hdfs hadoop 0 2021-11-20 19:07 /tmp
drwxr-xr-x - hdfs hadoop 0 2021-11-20 19:05 /user
drwxr-xr-x - hdfs hadoop 0 2021-11-20 19:05 /var
```

2. Creating a directory for our case study

Create directory command: ***hadoop fs -mkdir /user/hive/hivecasestudy***

Verifying the directory command: ***hadoop fs -ls /user/hive/***

```
[hadoop@ip-172-31-44-156 ~]$ hadoop fs -mkdir /casestudy /user/hive/hivecasestudy
[hadoop@ip-172-31-44-156 ~]$ hadoop fs -ls /user/hive/
Found 2 items
drwxr-xr-x - hadoop hadoop 0 2021-11-20 19:45 /user/hive/hivecasestudy
drwxrwxrwt - hdfs hadoop 0 2021-11-20 19:05 /user/hive/warehouse
[hadoop@ip-172-31-44-156 ~]$
```

From the above screenshot, we can see that the new directory is successfully created.

3. Load the data from S3 bucket to the HDFS

A. For our clickstream data, Month of Oct-2019.

Command:

hadoop distcp s3://hivecasestudy19/2019-Oct.csv User/hive/hivecasestudy/2019-Oct.csv

```
drwxrwxrwt - hdfs hadoop 0 2021-11-20 19:05 /user/hive/warehouse
[hadoop@ip-172-31-44-156 ~]$ hadoop distcp s3://hivecasestudy19/2019-Oct.csv /user/hive/hivecasestudy/2019-Oct.csv
21/11/20 19:50:11 INFO tools.DistCp: Input Options: DistCpOptions{atomicCommit=false, syncFolder=false, deleteMissing=false, ignoreFailures=false, overwrite=false, skipCRC=false, blocking=true, numListStatusThreads=0, maxMaps=20, mapBandwidth=100, sslConfigurationFile='null', copyStrategy='uniformsize', preserveStatus=[], preserveRawAttrs=false, atomicWorkPath=null, logPath=null, sourceFileListing=null, sourcePaths=[s3://hivecasestudy19/2019-Oct.csv], targetPath=/user/hive/hivecasestudy/2019-Oct.csv, targetPathExists=false, filtersFile='null'}
21/11/20 19:50:11 INFO client.RMProxy: Connecting to ResourceManager at ip-172-31-44-156.ec2.internal/172.31.4
```

B. Clickstream data, Month of Nov-2019.

Command:

hadoop distcp s3://hivecasestudy19/2019-Nov.csv /user/hive/hivecasestudy/2019-Nov.csv

```
[hadoop@ip-172-31-44-156 ~]$ hadoop distcp s3://hivecasestudy19/2019-Nov.csv /user/hive/hivecasestudy/2019-Nov.csv
21/11/20 19:53:39 INFO tools.DistCp: Input Options: DistCpOptions{atomicCommit=false, syncFolder=false, deleteMissi
ng=false, ignoreFailures=false, overwrite=false, skipCRC=false, blocking=true, numListStatusThreads=0, maxMaps=20,
mapBandwidth=100, sslConfigurationFile='null', copyStrategy='uniformsize', preserveStatus=[], preserveRawXattrs=fal
se, atomicWorkPath=null, logPath=null, sourceFileListing=null, sourcePaths=[s3://hivecasestudy19/2019-Nov.csv], tar
getPath=/user/hive/hivecasestudy/2019-Nov.csv, targetPathExists=false, filtersFile='null'}
21/11/20 19:53:40 INFO client.RMProxy: Connecting to ResourceManager at ip-172-31-44-156.ec2.internal/172.31.44.156
:8032
```

Verifying the loaded data in HDFS. Command: ***hadoop fs -ls /user/hive/hivecasestudy***

```
[hadoop@ip-172-31-44-156 ~]$ hadoop fs -ls /user/hive/hivecasestudy
Found 2 items
-rw-r--r-- 1 hadoop hadoop 545839412 2021-11-20 19:54 /user/hive/hivecasestudy/2019-Nov.csv
-rw-r--r-- 1 hadoop hadoop 482542278 2021-11-20 19:50 /user/hive/hivecasestudy/2019-Oct.csv
[hadoop@ip-172-31-44-156 ~]$
```

From the above screenshots, we can see that both of our files has been successfully loaded.

4. View the loaded data

Command: ***hadoop fs -cat /user/hive/hivecasestudy/2019-Oct.csv |head***

```
[hadoop@ip-172-31-44-156 ~]$ hadoop fs -cat /user/hive/hivecasestudy/2019-Oct.csv |head
event_time,event_type,product_id,category_id,category_code,brand,price,user_id,user_session
2019-10-01 00:00:00 UTC,car,5773203,1487580005134238553,,runail,2.62,463240011,26dd6e6e-4dac-4778-8d2c-92e149dab88
5
2019-10-01 00:00:03 UTC,car,5773353,1487580005134238553,,runail,2.62,463240011,26dd6e6e-4dac-4778-8d2c-92e149dab88
5
2019-10-01 00:00:07 UTC,car,5881589,2151191071051219817,,lovely,13.48,429681830,49e8d843-adf3-428b-a2c3-fe8bc6a307
c9
2019-10-01 00:00:07 UTC,car,5723490,1487580005134238553,,runail,2.62,463240011,26dd6e6e-4dac-4778-8d2c-92e149dab88
5
2019-10-01 00:00:15 UTC,car,5881449,1487580013522845895,,lovely,0.56,429681830,49e8d843-adf3-428b-a2c3-fe8bc6a307c
9
2019-10-01 00:00:16 UTC,car,5857269,1487580005134238553,,runail,2.62,430174032,73deale7-664e-43f4-8b30-d32b9d5af04
e
2019-10-01 00:00:19 UTC,car,5739055,1487580008246412266,,kapous,4.75,377667011,81326ac6-daa4-4f0a-b488-fd0956a7873
8
2019-10-01 00:00:24 UTC,car,5825598,1487580009445982239,,0.56,467916806,2f5b5546-b8cb-9ee7-7ecd-84276f8ef486
2019-10-01 00:00:25 UTC,car,5698989,1487580006317032337,,,1.27,385985999,d30965e8-1101-44ab-b45d-cc1bb9fae694
cat: Unable to write to output stream.
[hadoop@ip-172-31-44-156 ~]$
```

Command: ***hadoop fs -cat /user/hive/hivecasestudy/2019-Nov.csv |head***

```
cat: Unable to write to output stream.
[hadoop@ip-172-31-44-156 ~]$ hadoop fs -cat /user/hive/hivecasestudy/2019-Nov.csv |head
event_time,event_type,product_id,category_id,category_code,brand,price,user_id,user_session
2019-11-01 00:00:02 UTC,view,5802432,1487580009286598681,,,0.32,562076640,09fafd6c-6c99-46b1-834f-33527f4de241
2019-11-01 00:00:09 UTC,cart,5844397,1487580006317032337,,,2.38,553329724,2067216c-31b5-455d-alcc-af0575a34ffb
2019-11-01 00:00:10 UTC,view,5837166,1783999064103190764,,pnb,22.22,556138645,57ed222e-a54a-4907-9944-5a875c2d7f4f
2019-11-01 00:00:11 UTC,cart,5876812,1487580010100293687,,jessnail,3.16,564506666,186c1951-8052-4b37-adce-dd9644b1d5f7
2019-11-01 00:00:24 UTC,remove_from_cart,5826182,1487580007483048900,,,3.33,553329724,2067216c-31b5-455d-alcc-af0575a34ffb
2019-11-01 00:00:24 UTC,remove_from_cart,5826182,1487580007483048900,,,3.33,553329724,2067216c-31b5-455d-alcc-af0575a34ffb
2019-11-01 00:00:25 UTC,view,5856189,1487580009026551821,,runail,15.71,562076640,09fafd6c-6c99-46b1-834f-33527f4de241
2019-11-01 00:00:32 UTC,view,5837835,1933472286753424063,,,3.49,514649199,432a4e95-375c-4b40-bd36-0fc039e77580
2019-11-01 00:00:34 UTC,remove_from_cart,5870838,1487580007675986893,,milv,0.79,429913900,2f0bff3c-252f-4fe6-afcd-5d8a6a92839a
cat: Unable to write to output stream.
[hadoop@ip-172-31-44-156 ~]$
```

Create database

Launch Hive: Hadoop > Hive.

```
cat: Unable to write to output stream.
[hadoop@ip-172-31-44-156 ~]$ hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false
hive>
```

1. Creating a database

Command: ***Create database if not exists ecom;***

2. Use database created

Command : ***use ecom;***

3. Verifying the database created

Command: ***show databases;***

```
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false
hive> create database if not exists ecom;
OK
Time taken: 0.944 seconds
hive> use ecom;
OK
Time taken: 0.052 seconds
hive> show databases;
OK
default
ecom
Time taken: 0.173 seconds, Fetched: 2 row(s)
hive>
```

Create a table

1. Create a table using CSVSerde

Command:

CREATE EXTERNAL TABLE IF NOT EXISTS retail (event_time timestamp, event_type string, product_id string, category_id string, category_code string, brand string, price float,user_id bigint, user_session string) ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' STORED AS TEXTFILE LOCATION '/user/hive/hivecasestudy' tblproperties("skip.header.line .count"="1");

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS retail (event_time timestamp, event_type string, product_id string, category_id string, category_code string, brand string, price float,user_id bigint, user_session string) ROW FORMAT SERDE 'org.apach
e.hadoop.hive.serde2.OpenCSVSerde' STORED AS TEXTFILE LOCATION '/user/hive/hivecasestudy' tblproperties("skip.header.line .count"="1");
OK
Time taken: 0.371 seconds
```

2. Set the display for the header column
Command: **set hive.cli.print.header = true;**

```
Time taken: 0.371 seconds
hive> set hive.cli.print.header = true;
hive>
```

3. Verifying the table creation by checking the top 5 rows in the table.
Command: **select * from retail limit 5;**

```
hive> select * from retail limit 5;
OK
retail.event_time      retail.event_type      retail.product_id      retail.category_id      retail.category_code      retail
1.brand retail.price    retail.user_id      retail.user_session
2019-11-01 00:00:02 UTC view      5802432 1487580009286598681      0.32  562076640      09fafd6c-6c99-46b1-83
4f-33527f4de241
2019-11-01 00:00:09 UTC cart      5844397 1487580006317032337      2.38  553329724      2067216c-31b5-455d-a1
cc-af0575a34ffb
2019-11-01 00:00:10 UTC view      5837166 1783999064103190764      pnb  22.22  556138645      57ed222e-a54a-4907-99
44-5a875c2d7f4f
2019-11-01 00:00:11 UTC cart      5876812 1487580010100293687      jessnail  3.16  564506666      186c1951-8052
-4b37-adce-dd9644bd5f7
2019-11-01 00:00:24 UTC remove_from_cart      5826182 1487580007483048900      3.33  553329724      20672
16c-31b5-455d-a1cc-af0575a34ffb
Time taken: 3.9 seconds, Fetched: 5 row(s)
hive>
```

STEPS: APPLYING OPTIMIZATION TECHNIQUES (PARTITIONING & BUCKETING)

Enable Dynamic Partitioning

Command:

```
set hive.exec.dynamic.partition.mode = nonstrict;
set hive.exec.dynamic.partition = true;
```

Enable Bucketing

Command: **set hive.enforce.bucketing = true;**

```
hive> set hive.exec.dynamic.partition.mode = nonstrict;
hive> set hive.exec.dynamic.partition = true;
hive> set hive.enforce.bucketing = true;
hive>
```

1. Create an optimized table

Command:

```
CREATE TABLE IF NOT EXISTS dynpart_bucket_retail(event_time timestamp, product_id string,
category_id string, category_code string, brand string, price float, user_id bigint, user_session
string) PARTITIONED BY (event_type string) CLUSTERED BY (price) INTO 10 BUCKETS ROW
FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' STORED AS TEXTFILE
LOCATION '/user/hive/hivecasestudy' tblproperties('skip.header.line.count' = '1');
```

```
hive> set hive.enforce.bucketing = true;
hive> CREATE TABLE IF NOT EXISTS dynpart_bucket_retail(event_time timestamp, product_id string, category_id string, category_code
string, brand string, price float, user_id bigint, user_session string) PARTITIONED BY (event_type string) CLUSTERED BY (price) IN
TO 10 BUCKETS ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' STORED AS TEXTFILE LOCATION '/user/hive/hivecasestudy'
tblproperties('skip.header.line.count' = '1');
OK
Time taken: 0.071 seconds
```

For this table optimization based on the partitioning and bucketing, we have decided to create a partition on 'event_type' into 10 buckets and clustered by 'price'.

2. Verifying the table

Command: **show tables;**

```
hive> show tables;
OK
tab_name
dynpart_bucket_retail
retail
Time taken: 0.043 seconds, Fetched: 2 row(s)
```

As you can see, the table with partitions and bucketing named 'dynpart_bucket_retail' has been successfully created.

3. Insert the data into the optimized table

Command:

INSERT INTO TABLE dynpart_bucket_retail PARTITION (event_type) SELECT event_time, product_id, category_id, category_code, brand, price, user_id, user_session, event_type FROM retail;

```
hive>
> INSERT INTO TABLE dynpart_bucket_retail PARTITION (event_type) SELECT event_time, product_id, category_id, category_code, brand
, price, user_id, user_session, event_type FROM retail;
Query ID = hadoop_20211120201634_f087e19e-5e4a-43ac-b944-ac13ed4c2f8e
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1637435175999_0004)

-----
VERTICES      MODE           STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   2         2           0         0         0         0
Reducer 2 ..... container  SUCCEEDED   5         5           0         0         0         0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 163.69 s
-----
Loading data to table ecom.dynpart_bucket_retail partition (event_type=null)

Loaded : 4/4 partitions.
Time taken to load dynamic partitions: 1.026 seconds
Time taken for adding to write entity : 0.004 seconds
OK
event_time      product_id      category_id      category_code      brand      price      user_id user_session      event_type
Time taken: 175.334 seconds
hive>
```

4. Verifying the table created in Hadoop

Command: **hadoop fs -ls /user/hive/hivecasestudy**

```
[hadoop@ip-172-31-44-156 ~]$ hadoop fs -ls /user/hive/hivecasestudy
Found 6 items
-rw-r--r-- 1 hadoop hadoop 545839412 2021-11-20 19:54 /user/hive/hivecasestudy/2019-Nov.csv
-rw-r--r-- 1 hadoop hadoop 482542278 2021-11-20 19:50 /user/hive/hivecasestudy/2019-Oct.csv
drwxr-xr-x - hadoop hadoop 0 2021-11-20 20:19 /user/hive/hivecasestudy/event_type=cart
drwxr-xr-x - hadoop hadoop 0 2021-11-20 20:19 /user/hive/hivecasestudy/event_type=purchase
drwxr-xr-x - hadoop hadoop 0 2021-11-20 20:19 /user/hive/hivecasestudy/event_type=remove_from_cart
drwxr-xr-x - hadoop hadoop 0 2021-11-20 20:19 /user/hive/hivecasestudy/event_type=view
```

As you can see the partitioned files has been created and partitioned by 'event_type' which consists of cart, purchase, remove_from_cart and view. Let's explore further on the partitioning.

Command: ***hadoop fs -ls /user/hive/hivecasestudy/event_type=purchase***

```
[hadoop@ip-172-31-44-156 ~]$ hadoop fs -ls /user/hive/hivecasestudy/event_type=purchase
Found 10 items
-rwxr-xr-x 1 hadoop hadoop 6241877 2021-11-20 20:18 /user/hive/hivecasestudy/event_type=purchase/000000_0
-rwxr-xr-x 1 hadoop hadoop 7235640 2021-11-20 20:19 /user/hive/hivecasestudy/event_type=purchase/000001_0
-rwxr-xr-x 1 hadoop hadoop 7231471 2021-11-20 20:19 /user/hive/hivecasestudy/event_type=purchase/000002_0
-rwxr-xr-x 1 hadoop hadoop 7526313 2021-11-20 20:18 /user/hive/hivecasestudy/event_type=purchase/000003_0
-rwxr-xr-x 1 hadoop hadoop 7227979 2021-11-20 20:18 /user/hive/hivecasestudy/event_type=purchase/000004_0
-rwxr-xr-x 1 hadoop hadoop 7310389 2021-11-20 20:18 /user/hive/hivecasestudy/event_type=purchase/000005_0
-rwxr-xr-x 1 hadoop hadoop 8915123 2021-11-20 20:19 /user/hive/hivecasestudy/event_type=purchase/000006_0
-rwxr-xr-x 1 hadoop hadoop 5366094 2021-11-20 20:19 /user/hive/hivecasestudy/event_type=purchase/000007_0
-rwxr-xr-x 1 hadoop hadoop 6469070 2021-11-20 20:18 /user/hive/hivecasestudy/event_type=purchase/000008_0
-rwxr-xr-x 1 hadoop hadoop 8004214 2021-11-20 20:18 /user/hive/hivecasestudy/event_type=purchase/000009_0
[hadoop@ip-172-31-44-156 ~]$
```

For the 'event_type=purchase', there are exactly total of 10 buckets has been created.

Verifying the performance for the tables before and after optimization.

A. Retail table without optimization

Command: ***select * from retail limit 5;***

```
hive> select * from retail limit 5;
OK
2019-11-01 00:00:02 UTC view 5802432 1487580009286598681 0.32 562076640 09fafd6c-6c99-46b1-834f-3352
7f4de241
2019-11-01 00:00:09 UTC cart 5844397 1487580006317032337 2.38 553329724 2067216c-31b5-455d-a1cc-af05
75a34ffb
2019-11-01 00:00:10 UTC view 5837166 1783999064103190764 pnb 22.22 556138645 57ed222e-a54a-4907-9944-5a87
5c2d7f4f
2019-11-01 00:00:11 UTC cart 5876812 1487580010100293687 jessnail 3.16 564506666 186c1951-8052-4b37-a
dce-dd9644b1d5f7
2019-11-01 00:00:24 UTC remove_from_cart 5826182 1487580007483048900 3.33 553329724 2067216c-31b
5-455d-a1cc-af0575a34ffb
Time taken: 2.355 seconds, Fetched: 5 row(s)
hive>
```

B. Dynpart_bucket_retail table with optimization

Command: ***select * from dynpart_bucket_retail limit 5;***

```
hive> select * from dynpart_bucket_retail limit 5;
OK
2019-10-08 09:19:19 UTC 89350 1487580011652186237 runail 1.27 232701853 3f1469f5-d926-44ce-a9f6-dff5ae276c9c
cart
2019-10-10 05:29:47 UTC 5866208 1487580013841613016 concept 3.16 493381333 535bb6b7-08f4-4021-ac66-b340178f7a37
cart
2019-10-08 12:25:50 UTC 5821183 1487580007717929935 1.27 546703849 3daf4d64-5ffa-46cc-827b-59760ebd819b
cart
2019-10-10 08:19:06 UTC 5848901 1487580007675986893 bpw.style 1.27 439370683 9aeb4d9a-1bed-4f42-b12d-88be
1148d3a9 cart
2019-10-09 18:32:50 UTC 5869152 1487580005268456287 cosmoprofi 7.94 558533352 cfde0f74-8705-4a2f-ba83-a5b9
9581c294 cart
Time taken: 0.204 seconds, Fetched: 5 row(s)
hive>
```

Insights:

1. Before optimization for the 'retail table' it tooks 2.355 seconds to retrieve the data from the query, however with the dynpart_bucket_table (after optimization), it tooks only 0.204 seconds to read the data. Difference of 2.151 seconds, it's too early to make any assumptions, lets look into further queries.

QUESTIONS

Question 1: Find the total revenue generated due to purchases made in October.

Without Optimization (Retail Table)

Command:

select sum(price) as october_revenue from retail where month(event_time) = '10' AND event_type = 'purchase';

```
hive> select sum(price) as october_revenue from retail where month(event_time) = '10' AND event_type = 'purchase';
Query ID = hadoop_20211120203314_9a06399d-ea55-4e36-b12c-ebb8b6bd86d6
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1637435175999_0006)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    5         5         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 126.05 s
-----
OK
1211538.4299997438
Time taken: 135.866 seconds, Fetched: 1 row(s)
hive>
```

With Optimization (Dynpart_bucket_retail Table)

Command:

select sum(price) as october_revenue from dynpart_bucket_retail where month(event_time) = '10' AND event_type = 'purchase';

```
hive> select sum(price) as october_revenue from dynpart_bucket_retail where month(event_time) = '10' AND event_type = 'purchase';
Query ID = hadoop_20211120203624_af2630a7-579b-4688-af0c-2ca6dcd88f91
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1637435175999_0006)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    3         3         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 23.25 s
-----
OK
1211532.4500002791
Time taken: 24.25 seconds, Fetched: 1 row(s)
hive>
```

Insights:

1. In this query, we are required to find the total revenue generated on October 2019 based on the Purchases.
2. The time taken to read the data for the 'retail table (without optimization)' is 135.866 seconds, while for the 'dynpart_bucket_retail table (with optimization)' is 24.25 seconds, which is 5x lesser than the retail table.
3. To answer the above questions, there are total of 1211538.43 +/- revenue generate on October on Purchases.
4. There are huge differences in the time taken to retrieve the data from both tables, in this query, the optimization table performed faster than 'retail table' which is without optimization.

Question 2: Write a query to yield the total sum of purchases per month in a single output.

Without Optimization (Retail Table)

Command:

select month(event_time) as month, count(event_type) as total_purchases from retail where event_type = 'purchase' group by month(event_time);

```
hive> select month(event_time) as month, count(event_type) as total_purchases from retail where event_type = 'purchase' group by month(event_time);
Query ID = hadoop_20211120204414_244640f9-cb71-4cf9-a021-12037757d42e
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1637435175999_0007)

-----
VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    5         5         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    3         3         0         0         0         0
-----
VERTICES: 02/02  [=====>>>] 100%  ELAPSED TIME: 102.25 s
-----
OK
10      245624
11      322417
Time taken: 102.853 seconds, Fetched: 2 row(s)
hive>
```

With Optimization (Dynpart_bucket_retail Table)

Command:

select month(event_time) as month, count(event_type) as total_purchases from dynpart_bucket_retail where event_type = 'purchase' group by month(event_time);

```
hive> select month(event_time) as month, count(event_type) as total_purchases from dynpart_bucket_retail where event_type = 'purchase' group by month(event_time);
Query ID = hadoop_20211120204702_b836adb6-2635-4453-b78a-053b18fe0890
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1637435175999_0007)

-----
VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    3         3         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 02/02  [=====>>>] 100%  ELAPSED TIME: 23.28 s
-----
OK
10      245619
11      322412
Time taken: 24.022 seconds, Fetched: 2 row(s)
hive>
```

Insights:

1. In this next question, we are required to find the total sum for the purchases for the month of October and November.
2. The Optimized table again prove the faster query with only 24.02 seconds for the time taken and the ratil table took 102.85 seconds to retrieve the data.
3. For the month of October, the total purchases is 245624 and 322417 for the month of November. There are a significant improvement on the purchase value which is around 30%.

Question 3: Write a query to find the change in revenue generated due to purchases from October to November.

Without Optimization (Retail Table)

Command:

select (sum(case when month(event_time)=11 then price else 0 end) - sum(case when month(event_time)=10 then price else 0 end)) as change_in_rev from retail where event_type = 'purchase' and month(event_time) in ('10','11');

```
hive> select (sum(case when month(event_time)=11 then price else 0 end) - sum(case when month(event_time)=10 then price else 0 end)) as change_in_rev from retail where event_type = 'purchase' and month(event_time) in ('10','11');
Query ID = hadoop_20211120205309_820e6dba-af39-438b-8830-a126b905d03f
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1637435175999_0008)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    5          5          0          0          0          0
Reducer 2 ..... container  SUCCEEDED    1          1          0          0          0          0
-----
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 102.63 s
-----
OK
319478.4700003781
Time taken: 112.009 seconds, Fetched: 1 row(s)
```

With Optimization (Dyngpart_bucket_retail Table)

Command:

select (sum(case when month(event_time)=11 then price else 0 end) - sum(case when month(event_time)=10 then price else 0 end)) as change_in_rev from dynpart_bucket_retail where event_type = 'purchase' and month(event_time) in ('10','11');

```
hive> select (sum(case when month(event_time)=11 then price else 0 end) - sum(case when month(event_time)=10 then price else 0 end)) as change_in_rev from dynpart_bucket_retail where event_type = 'purchase' and month(event_time) in ('10','11');
Query ID = hadoop_20211120205605_3d20ec40-d72c-436c-9fd2-de05c52d2018
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1637435175999_0008)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    3          3          0          0          0          0
Reducer 2 ..... container  SUCCEEDED    1          1          0          0          0          0
-----
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 25.44 s
-----
OK
319437.7899997565
Time taken: 26.109 seconds, Fetched: 1 row(s)
```

Insights:

1. In this next question, we will find out on the differences in the revenue generated based on purchases for the month of October and November.
2. The difference of revenue generated between the month of October and November 2019 are 319478.47 +/-.
3. It took 112.009 second for the Retail table to read the data from the above query, however only 26.109 seconds time taken from the optimization table. Again, the Optimized table performed faster in this analysis. Hence, for the next remaining questions, we will be using the optimized tables to run the queries.

Question 4: Find distinct categories of products. Categories with null category code can be ignored.

With Optimization (Dynpart_bucket_retail Table)

Command:

select distinct split(category_code,'\\')[0] as distinct_category from dynpart_bucket_retail where category_code != '';

```
hive> select distinct split(category_code,'\\')[0] as distinct_category from dynpart_bucket_retail where category_code != '';
Query ID = hadoop_20211120205910_a541c0d4-a105-4d6f-ba75-a28001eae34c
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1637435175999_0008)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	6	container	SUCCEEDED	6	6	0	0	0	0
Reducer 2	5	container	SUCCEEDED	5	5	0	0	0	0

```
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 64.51 s
OK
furniture
appliances
accessories
apparel
sport
stationery
Time taken: 65.263 seconds, Fetched: 6 row(s)
```

Insights:

1. The time take to retrieve the data to execute the query is 65.263 seconds.
2. This question required us to find out the categories present in the data, and there are total of 6 categories available, which are:
 - Furniture
 - Appliances
 - Accessories
 - Sport
 - Apparel
 - Stationery

Question 5: Find the total number of products available under each category

With Optimization (Dynpart_bucket_retail Table)

Command:

```
select split(category_code,'\\')[0] as category, count(product_id) as num_of_products from  
dynpart_bucket_retail where category_code != "" group by split(category_code,'\\')[0] order by  
num_of_products desc;
```

```
hive> select split(category_code,'\\')[0] as category, count(product_id) as num_of_products from dynpart_bucket_retail where category_code != ''  
roup by split(category_code,'\\')[0] order by num_of_products desc;  
Query ID = hadoop_20211120210339_d5769307-9a16-462c-9855-dec7634ef342  
Total jobs = 1  
Launching Job 1 out of 1  
Status: Running (Executing on YARN cluster with App id application_1637435175999_0008)  
  
-----  
VERTICES      MODE           STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED  
-----  
Map 1 ..... container  SUCCEEDED    6         6          0         0         0         0  
Reducer 2 ..... container  SUCCEEDED    5         5          0         0         0         0  
Reducer 3 ..... container  SUCCEEDED    1         1          0         0         0         0  
-----  
VERTICES: 03/03 [=====>>>] 100% ELAPSED TIME: 64.70 s  
-----  
OK  
appliances      61736  
stationery      26722  
furniture       23604  
apparel 18232  
accessories     12928  
sport          2  
Time taken: 65.36 seconds, Fetched: 6 row(s)
```

Insights:

1. Based from the categories listed from the previous question, we are required to find the total number of products in each of the 6 categories.
2. To execute this query, it took 65.36 seconds.
3. And the output are as follows:

No.	Category	Number of Products
1.	Appliances	61736
2.	Stationery	26722
3.	Furniture	23604
4.	Apparel	18232
5.	Accessories	12928
6.	Sports	2

Appliances lead the total number of products with 61736, followed closely by Stationery and Furniture. Surprisingly Sports only consists of 2 products.

Question 6: Which brand had the maximum sales in October and November combined?

With Optimization (Dynpart_bucket_retail Table)

Command:

WITH total_sales_summary AS(

select brand, round((sum(case when month(event_time)=10 then price else 0 end) + sum(case when month(event_time)=11 then price else 0 end)),2) as total_sales from dynpart_bucket_retail where event_type = 'purchase' and month(event_time) in ('10','11') and brand != '' group by brand)

select brand, total_sales from total_sales_summary order by total_sales desc limit 1;

```
hive> WITH total_sales_summary AS(
  > select brand, round((sum(case when month(event_time)=10 then price else 0 end) + sum(case when month(event_time)=11 then price else 0 end)),2)
  > ) as total_sales
  > from dynpart_bucket_retail
  > where event_type = 'purchase' and month(event_time) in ('10','11') and brand != '' group by brand)
  > select brand, total_sales from total_sales_summary order by total_sales desc limit 1;
Query ID = hadoop_20211120211319_4a9165b4-51c1-4637-9915-bd8683a38ea5
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1637435175999_0009)

-----
VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   3         3         0         0         0         0
Reducer 2 ..... container  SUCCEEDED   1         1         0         0         0         0
Reducer 3 ..... container  SUCCEEDED   1         1         0         0         0         0
-----
VERTICES: 03/03  [=====>>] 100% ELAPSED TIME: 27.14 s
-----
OK
runail 148292.46
Time taken: 35.324 seconds, Fetched: 1 row(s)
hive>
```

Insights:

1. In this question, we are required to find the name of the brand which had the total number of sales for BOTH month of October and November (combine).
2. The time taken to retrieve the data from the above query is only 35.324 seconds.
3. It turns out that brand Runail possessed the highest sales for the month of October AND November.

Question 7: Which brands increased their sales from October to November?

With Optimization (Dynpart_bucket_retail Table)

Command:

WITH brand_sales_summary AS(

select brand, round(sum(case when month(event_time)=10 then price else 0 end),2) as sales_october,round(sum(case when month(event_time)=11 then price else 0 end),2) as sales_november from dynpart_bucket_retail where event_type = 'purchase' and month(event_time) in ('10','11') and brand != '' group by brand)

select brand, sales_october, sales_november, round((sales_november - sales_october),2)as sales_differences from brand_sales_summary where sales_november - sales_october > 0 order by sales_differences desc;

```
hive> WITH brand_sales_summary AS(
  > select brand, round(sum(case when month(event_time)=10 then price else 0 end),2) as sales_october,round(sum(case when month(event_time)=11 then price else 0 end),2) as sales_november from dynpart_bucket_retail where event_type = 'purchase' and month(event_time) in ('10','11') and brand != '' group by brand)
  > select brand, sales_october, sales_november, round((sales_november - sales_october),2)as sales_differences from brand_sales_summary where sales_november - sales_october > 0 order by sales_differences desc;
Query ID = hadoop_20211120212218_9daf643a-84f8-4e88-892f-70bb81764058
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1637435175999_0010)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	3	3	0	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0	0
Reducer 3	container	SUCCEEDED	1	1	0	0	0	0	0

VERTICES: 03/03 [=====] 100% ELAPSED TIME: 27.78 s

OK

grattol	35445.54	71472.71	36027.17
uno	35302.03	51039.75	15737.72
lianail	5892.84	16394.24	10501.4
ingarden	23161.39	33566.21	10404.82
strong	29196.63	38671.27	9474.64
jessnail	26287.84	33345.23	7057.39
cosmoprofi	8322.81	14536.99	6214.18
polarus	6013.72	11371.93	5358.21
runail	71537.77	76754.69	5216.92
freedecor	3421.78	7671.8	4250.02
staleks	8519.73	11875.61	3355.88
bpw.style	11572.15	14837.44	3265.29
lovely	8704.38	11939.06	3234.68
marathon	7280.75	10273.1	2992.35
haruyama	9390.69	12352.91	2962.22

yoko	8756.91	11707.88	2950.97
italwax	21940.24	24799.37	2859.13
benovy	409.62	3259.97	2850.35
kaypro	881.34	3268.7	2387.36
estel	21756.75	24142.67	2385.92
concept	11032.14	13380.4	2348.26
kapous	11927.16	14093.08	2165.92
f.o.x	6624.23	8577.28	1953.05
masura	31266.08	33058.47	1792.39
milv	3904.94	5642.01	1737.07
beautix	10493.95	12222.95	1729.0
artex	2730.64	4327.25	1596.61
domix	10472.05	12009.17	1537.12
shik	3341.2	4839.72	1498.52
smart	4457.26	5902.14	1444.88
roubloff	3491.36	4913.77	1422.41
levrana	2243.56	3664.1	1420.54
oniq	8425.41	9841.65	1416.24
irisk	45591.96	46946.04	1354.08
severina	4775.88	6120.48	1344.6
joico	705.52	2015.1	1309.58
zeitun	708.66	2009.63	1300.97
beauty-free	554.17	1782.86	1228.69
swarovski	1887.93	3043.16	1155.23
de.lux	1659.7	2775.51	1115.81
metzger	5373.45	6457.16	1083.71
markell	1768.75	2834.43	1065.68
sanoto	157.14	1209.68	1052.54
nagaraku	4369.74	5327.68	957.94
ecolab	262.85	1214.3	951.45
art-visage	2092.71	2997.8	905.09
levissime	2227.5	3085.31	857.81
missha	1293.83	2150.28	856.45
solomeya	1899.7	2685.8	786.1
rosi	3077.04	3841.56	764.52
refectocil	2716.18	3475.58	759.4
kaaral	4412.43	5086.07	673.64
kosmekka	1181.44	1813.37	631.93

matreshka	0.0	182.67	182.67
chi	358.94	538.61	179.67
cristalinas		427.63	584.95
farmona	1692.46	1843.43	150.97
latinoil		249.52	384.59
maskin	158.04	293.07	135.03
elizavecca		70.53	204.3
nefertiti		233.52	366.64
finish	98.38	230.38	132.0
igrobeauty		513.66	645.07
dizao	819.13	945.51	126.38
osmo	645.58	762.31	116.73
batiste	772.4	874.17	101.77
carmex	145.08	243.36	98.28
eos	54.34	152.61	98.27
depilflax		2707.07	2803.78
enjoy	41.35	136.57	95.22
kerasys	430.91	525.2	94.29
aura	83.95	177.51	93.56
plazan	101.37	194.01	92.64
koelf	422.73	507.29	84.56
nirvel	163.04	234.33	71.29
konad	739.83	810.67	70.84
egomania		77.47	146.04
cutrin	299.37	367.62	68.25
laboratorium		246.5	312.52
inm	288.02	351.21	63.19
dewal	0.0	61.29	61.29
marutaka-foot		49.22	109.33
kares	0.0	59.45	59.45
profhenna		679.23	736.85
koelcia	55.5	112.75	57.25
balbcare		155.33	212.38
elskin	251.09	307.65	56.56
foamie	35.04	80.49	45.45
ladykin	125.65	170.57	44.92
likato	296.06	340.97	44.91
mavala	409.04	446.32	37.28

kinetics	6334.25	6945.26	611.01
browxenna		14331.37	14916.73
airnails		5118.9	5691.52
uskusi	5142.27	5690.31	548.04
coifin	903.0	1428.49	525.49
s.care	412.68	913.07	500.39
limoni	1308.9	1796.6	487.7
matrix	3243.25	3726.74	483.49
gehwol	1089.07	1557.68	468.61
greymy	29.21	489.49	460.28
bioaqua	942.89	1398.12	455.23
farmavita		837.37	1291.97
sophin	1067.86	1515.52	447.66
yu-r	271.41	673.71	402.3
kiss	421.55	817.33	395.78
naomi	0.0	389.0	389.0
lador	2083.61	2471.53	387.92
ellips	245.85	606.04	360.19
jas	3318.96	3657.43	338.47
lowence	242.84	567.75	324.91
nitriile	847.28	1162.68	315.4
shary	871.96	1176.49	304.53
kims	330.04	632.04	302.0
happyfons		801.92	1091.59
kocostar		310.85	594.93
insight	1443.7	1721.96	278.26
candy	534.96	799.38	264.42
bluesky	10307.24	10565.53	258.29
beauugreen		511.51	768.35
protokeratin		201.25	456.79
trind	298.07	542.96	244.89
entity	479.71	719.26	239.55
skinlite		651.94	890.45
provoc	827.99	1063.82	235.83
fedua	52.38	263.81	211.43
ecocraft		41.16	241.95
keen	236.35	435.62	199.27
mane	66.79	260.26	193.47
freshbubble		318.7	502.34

vilenta	197.6	231.21	33.61
beautyblender		78.74	109.41
biore	60.65	90.31	29.66
brly	902.38	931.09	28.71
estelare		444.81	471.87
profepil		93.36	118.02
blizx	38.95	63.4	24.45
binacil	0.0	24.26	24.26
godefroy		401.22	425.12
glysolid		69.73	91.59
veraclara		50.11	71.21
juno	0.0	21.08	21.08
kamill	63.01	81.49	18.48
treaclemoon		163.37	181.49
supertan		50.37	66.51
barbie	0.0	12.39	12.39
deoproce		316.84	329.17
rasyan	18.8	28.94	10.14
fly	17.14	27.17	10.03
tertio	236.16	245.8	9.64
jaguar	1102.11	1110.65	8.54
soleo	204.2	212.53	8.33
neoleor	43.41	51.7	8.29
moyou	5.71	10.28	4.57
bodyton	1376.34	1380.64	4.3
skinity	8.88	12.44	3.56
hellorganic		0.0	3.1
grace	100.92	102.61	1.69
cosima	20.23	20.93	0.7
ovale	2.54	3.1	0.56
Time taken: 36.658 seconds, Fetched: 160 row(s)			

Insights:

1. To answer this question, we need to find the name of the brand which has increased their sales from the month of October and November.
2. It took only 36.658 seconds to retrieve the data from the Optimized table.
3. From the output, we can see that there are the total of 160 rows, which means, there are the total of 160 brands which has increased their sales in the subsequent month.
4. Grattol if the leading brand with the total increment of 2x which is 36027.17 increment in both of the month, followed closely by Uno brand with 15737.32 and Lianail with 10501.40. The least brand which has the lowest differences in Sales for the month of October and November are Ovale with 0.56, Cosima 0.70 and Grace with 1.69.

Question 8: Your company wants to reward the top 10 users of its websites with a golden customer plan. Write a query to generate a list of top 10 users who spend the most.

With Optimization (Dynpart_bucket_retail Table)

Command:

select user_id, round(sum(price),2) as amount_spends from dynpart_bucket_retail where event_type = 'purchase' group by user_id order by amount_spends desc limit 10;

```
hive> select user_id, round(sum(price),2) as amount_spends from dynpart_bucket_retail where event_type = 'purchase' group by user_id order by amount_spends desc limit 10;
Query ID = hadoop_20211120212904_35da6421-cbf5-4130-a9a1-72b47e2f4501
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1637435175999_0011)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	3	3	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0
Reducer 3	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 03/03 [=====]>>] 100% ELAPSED TIME: 26.01 s
OK
557790271      2715.87
150318419      1645.97
562167663      1352.85
531900924      1329.45
557850743      1295.48
522130011      1185.39
561592095      1109.7
431950134      1097.59
566576008      1056.36
521347209      1040.91
Time taken: 34.113 seconds, Fetched: 10 row(s)
```

Insights:

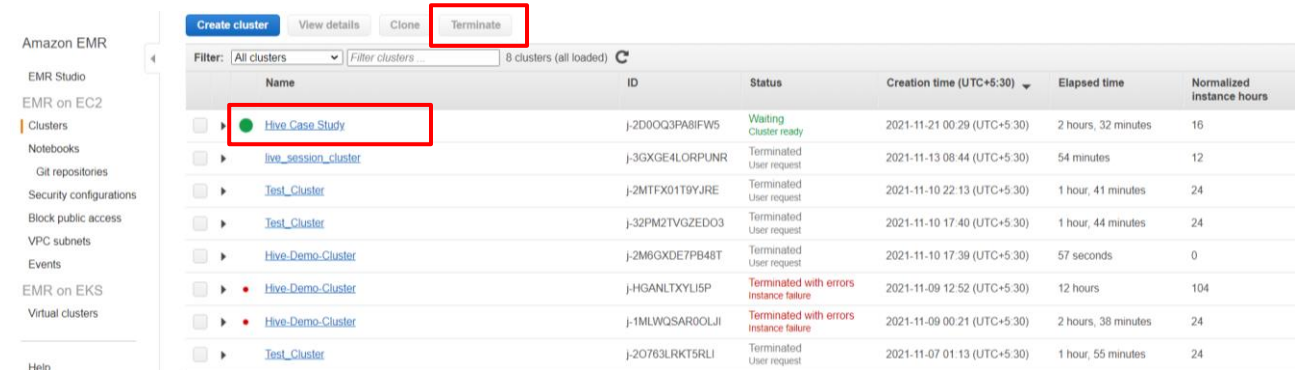
1. We need to find out the Top 10 users who spend the most on the website, which will won the Golden Customer Plan from our company.
2. The time taken for this query to be completed is 34.113 seconds.
3. Based from the output, we can see that the Top user spend 2715.87, followed by 1645.97 and 1352.85.

We have completed all of the questions, which needs to be run on the system, and we can conclude that Partitioning and Bucketing helps in increasing the performance by providing faster analysis on the query despite the huge amount of data provided. And it is efficient and convenient, as well as easy to code despite the volume of data loaded.

STEPS: TERMINATING EMR CLUSTER

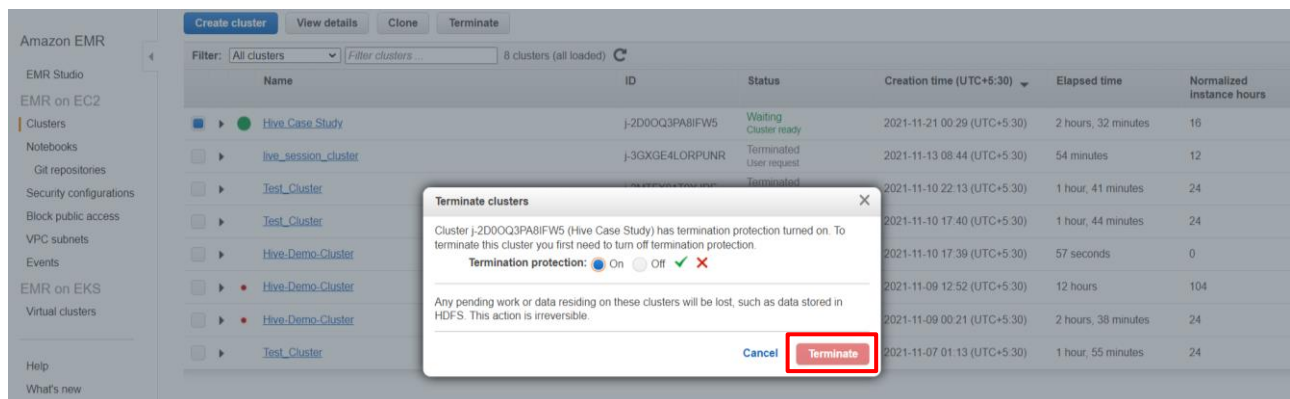
It is very important to terminate the cluster, because of the charges that will be charged into our account. Thus, we followed the below procedure to complete the above process:

On AWS Services, type EMR and go back to the main page of our Cluster.



Tick on our cluster, in this case 'Hive_Case_Study' cluster > click on Terminate.

A screen will pop up, we have unticked the 'On' and switch to 'Off' to allow the termination procedure to be completed, then click on 'Terminate'.



Our cluster and both of our Master and Core node has been terminated and with that we have completed ur hive Analysis on ClickStream dataset successfully.

